

You are currently looking at **version 1.2** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.

## Assignment 3 - Evaluation

In this assignment you will train several models and evaluate how effectively they predict instances of fraud using data based on [this dataset from Kaggle](#).

Each row in `fraud_data.csv` corresponds to a credit card transaction. Features include confidential variables `V1` through `V28` as well as `Amount` which is the amount of the transaction.

The target is stored in the `class` column, where a value of 1 corresponds to an instance of fraud and 0 corresponds to an instance of not fraud.

In [1]:

```
import numpy as np
import pandas as pd
```

### Question 1

Import the data from `fraud_data.csv`. What percentage of the observations in the dataset are instances of fraud?

*This function should return a float between 0 and 1.*

In [2]:

```
def answer_one():
    df = pd.read_csv('fraud_data.csv')
    y = df.iloc[:, -1]

    percent_fraud = len(y[y==1]) / len(y)

    return percent_fraud
#answer_one()
#Function answer_one was answered incorrectly, 0.165 points were not awarded.
```

In [3]:

```
# Use X_train, X_test, y_train, y_test for all of the following questions
from sklearn.model_selection import train_test_split

df = pd.read_csv('fraud_data.csv')

X = df.iloc[:, :-1]
y = df.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

### Question 2

Using `X_train`, `X_test`, `y_train`, and `y_test` (as defined above), train a dummy classifier that classifies everything as the majority class of the training data. What is the accuracy of this classifier? What is the recall?

*This function should return a tuple with two floats, i.e. (accuracy score, recall score).*

In [4]:

```
def answer_two():
    from sklearn.dummy import DummyClassifier
```

```

from sklearn.metrics import accuracy_score, recall_score
dummy_majority = DummyClassifier(strategy = 'most_frequent').fit(X_train, y_train)
y_dummy_predictions = dummy_majority.predict(X_test)
accuracy_score=accuracy_score(y_test, y_dummy_predictions)
recall_score=recall_score(y_test, y_dummy_predictions)
return (accuracy_score, recall_score)
#answer_two()

```

### Question 3

Using `X_train`, `X_test`, `y_train`, `y_test` (as defined above), train a SVC classifier using the default parameters. What is the accuracy, recall, and precision of this classifier?

*This function should return a tuple with three floats, i.e. (accuracy score, recall score, precision score).*

In [5]:

```

def answer_three():
    from sklearn.metrics import accuracy_score, recall_score, precision_score
    from sklearn.svm import SVC
    svm = SVC().fit(X_train, y_train)
    svm_predicted = svm.predict(X_test)
    accuracy_score=accuracy_score(y_test, svm_predicted)
    recall_score=recall_score(y_test, svm_predicted)
    precision_score=precision_score(y_test, svm_predicted)
    return (accuracy_score, recall_score, precision_score)
#answer_three()

```

### Question 4

Using the SVC classifier with parameters `{'C': 1e9, 'gamma': 1e-07}`, what is the confusion matrix when using a threshold of -220 on the decision function. Use `X_test` and `y_test`.

*This function should return a confusion matrix, a 2x2 numpy array with 4 integers.*

In [6]:

```

def answer_four():
    from sklearn.metrics import confusion_matrix
    from sklearn.svm import SVC
    svm = SVC(C=1e9, gamma=1e-07)
    y_scores_lr = svm.fit(X_train, y_train).decision_function(X_test)
    y_score_list = list(zip(y_test, y_scores_lr))
    y_predict_threshold=[]
    #for loop where every time the decision function "score" is greater than -220 make it a 1, else
    e make it a 0
    for i in y_score_list:
        if i[1]>-220:
            y_predict_threshold.append(1)
        else:
            y_predict_threshold.append(0)
    confusion = confusion_matrix(y_test, y_predict_threshold)
    return confusion
#answer_four()

```

### Question 5

Train a logistic regression classifier with default parameters using `X_train` and `y_train`.

For the logistic regression classifier, create a precision recall curve and a roc curve using `y_test` and the probability estimates for `X_test` (probability it is fraud).

Looking at the precision recall curve, what is the recall when the precision is 0.75?

Looking at the roc curve, what is the true positive rate when the false positive rate is 0.16?

*This function should return a tuple with two floats, i.e. (recall, true positive rate).*

In [7]:

In [ / ]:

```
def answer_five():
    from sklearn.linear_model import LogisticRegression
    #from sklearn.metrics import precision_recall_curve
    from sklearn.metrics import roc_curve, auc
    import matplotlib.pyplot as plt
    lr = LogisticRegression()
    #use decision_function, not predict
    y_score_lr = lr.fit(X_train, y_train).decision_function(X_test)
    #fpr_lr, tpr_lr, _ = roc_curve(y_test, y_score_lr)
    #roc_auc_lr = auc(fpr_lr, tpr_lr)

    #after having done the plot and eye-balled it
    recall=0.8
    #after having done the plot and eye-balled it
    true_pos=0.9
    return (recall, true_pos)

#answer_five()
```

## Question 6

Perform a grid search over the parameters listed below for a Logistic Regression classifier, using recall for scoring and the default 3-fold cross validation.

'penalty': ['l1', 'l2']

'C':[0.01, 0.1, 1, 10, 100]

From .cv\_results\_, create an array of the mean test scores of each parameter combination. i.e.

	l1	l2
0.01	?	?
0.1	?	?
1	?	?
10	?	?
100	?	?

This function should return a 5 by 2 numpy array with 10 floats.

Note: do not return a DataFrame, just the values denoted by '?' above in a numpy array. You might need to reshape your raw result to meet the format we are looking for.

In [12]:

```
def answer_six():
    from sklearn.model_selection import GridSearchCV
    from sklearn.linear_model import LogisticRegression
    grid_values = {'C':[0.01, 0.1, 1, 10, 100], 'penalty': ['l1', 'l2']}
    lr=LogisticRegression()
    grid_lr_recall = GridSearchCV(lr, param_grid = grid_values, scoring = 'recall')
    grid_lr_recall.fit(X_train, y_train)
    #cv_results_ returns a dictionary from 'mean_test_score' key get values
    #for l1 start at index 0 and get every other one (that's how they are arranged according to .c
    v_results_['param_penalty'])
    l1=grid_lr_recall.cv_results_['mean_test_score'][0::2]
    #reshape it into the preferred shape
    l1=np.array(l1).reshape(5,1)
    #for l2 start at index 1 and get every other one
    l2=grid_lr_recall.cv_results_['mean_test_score'][1::2]
    #reshape it into the preferred shape
    l2=np.array(l2).reshape(5,1)
    return np.hstack((l1, l2))

#answer_six()
#Function answer_five was answered incorrectly, 0.167 points were not awarded. The code for this p
roblem is not compiling.
```

In [10]:

```
# Use the following function to help visualize results from the grid search
def GridSearch_Heatmap(scores):
    %matplotlib notebook
    import seaborn as sns
    import matplotlib.pyplot as plt
    plt.figure()
    sns.heatmap(scores.reshape(5,2), xticklabels=['l1','l2'], yticklabels=[0.01, 0.1, 1, 10, 100])
    plt.yticks(rotation=0);

#GridSearch_Heatmap(answer_six())
```